


In-distribution training games ($s_t \rightarrow \text{NCA-WM} \rightarrow \hat{s}_{t+1}$; engine s_{t+1} supervises the loss)



Held-out (OOD) games (only s_0 fed in; $\hat{s}_{t+1}$ rolls forward autoregressively, no supervision)

